



Predictive Modeling for Direct Marketing

Bank Case

September 2019



Contents

Executive Summary

Analytics Goal and Approach

Appendix

Importance of direct marketing decision support through predictive customer response modeling is recognized

Background

- A bank is marketing Certificates of Deposit (CD) to a pool of 1,000 potential customers. The bank estimates that 13% of them are likely to buy a CD if they are contacted.
- The cost of each contact is \$10. The NPV to the bank of a customer buying a CD is \$50. Clearly, it is not profitable to contact all the potential customers.
- It has an unbalanced data set* (bank.csv) of some customers and whether they bought the CD or not. The data set has 16 features and a class variable for a total of 17 columns (e.g., age, job, marital, education, default, housing, balance, loan, campaign). It has 4,521 observations.



Request for Analytics Assistance

- To maximize expected profit, the bank is expected to employ a predictive model to predict whether or not a potential customer is a likely buyer.
- The data set will be used here to build the predictive model. After unbalanced data set is adjusted and decision tree is calibrated using complexity parameter and optimal cutoff, the expected profit is maximized at \$2261.017.
- Given the goal of driving future profit by better predicting which of past customers are best prospects to purchase CDs, the implementation of analytics team should focus on customer data organization, model performance perfection, and results monitoring.

*See dataset description in Appendix.



Contents

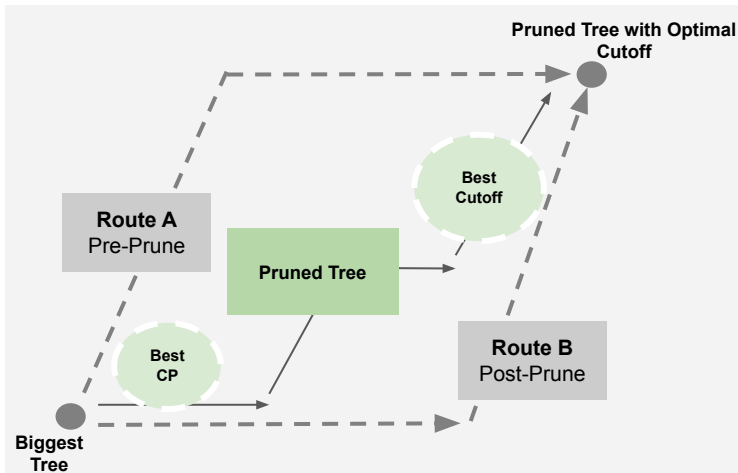
Executive Summary

Analytics Goal and Approach

Appendix

Decision tree technique will be applied for target selection, reaching objective of direct marketing to maximize expected profit

A Decision Tree Pruned Using Complexity Parameter (CP) and Optimal Cutoff



- Route A (Pre-Prune): `fit = rpart(churn ~ ., data=b.train, control=rpart.control(xval=10, cp=bestcp))`
- Route B (Post-Prune): `fit = prune.rpart(fit, cp=bestcp)`
- Two routes produce similar results, thus in code post-pruned way is used as example.
- In the next slide, how to get optimal tree by finding best cp and best cutoff in a balanced dataset will be discussed.

Maximum Profit by Building A Predictive Model with Good Quality

	Positive	Negative	
Positive	$130 * (1 - \text{FNR}) * \40	$- 870 * \text{FPR} * \$10$	Predicted
Negative	\$0	\$0	
	Actual		

Expected Profit of Using Predictive Model to Contact Customers Predicted Positive:

$$130 * (1 - \text{FNR}) * 40 - 870 * \text{FPR} * 10$$

$$= 5200 - 5200 * \text{FNR}(\text{beta}) - 8700 * \text{FPR}(\text{alpha})$$

Analytics Design & Process - Workstream Highlights*

1

Build a complex tree

2

Post prune with best complexity parameter

3

Create a balanced dataset

4

Retrain model using balanced dataset

5

Find Optimal Threshold for Cutoff

```
bestcp =
fit$cpstable[which.min(fit$cpstable[, "xerror"]),
"CP"]
# lowest error occurs at cp = 0.01186944
```

```
# post-pruning
fit.post = prune.rpart(fit, cp=bestcp)
nrow(fit.post$frame)
# the pruned tree have 17 nodes
```

```
# compute the confusion matrices in test set
cm = confusionMatrix(table(predict(fit.post,
b.test, type="class"), b.test$y),
positive='yes')
```

```
profit = 5200 - 5200 *
(1-cm$byClass["Sensitivity"])[[1]] - 8700 *
(1-cm$byClass["Specificity"])[[1]]
# profit = $1600.222
```

```
# splitting the training
dataset
b.train.yes = subset(b.train,
y == 'yes')
b.train.no = subset(b.train, y
== 'no')
```

```
# take a subsample from
b.train.no with the same
number of observations as
b.train.yes.
set.seed(234)
no =
sample(1:nrow(b.train.no),nr
ow(b.train.yes))
b.train.no = b.train.no[no,]
```

```
# Combine b.train.yes and
subsample of b.train.no to
make a new balanced
training data frame
b.bal = rbind(b.train.yes,
b.train.no)
```

```
# bal.fit has 95 nodes
```

```
# lowest error occurs at
bal.bestcp = 0.007418398
```

```
# post-pruned tree has
bal.fit.post 19 nodes
```

```
# profit = $2186.02
```

```
library(rpart)
# xval = 10 : cross-sample with 10 folds to determine error rate at each node
# minsplitt = 10 : min number of observations to attempt split
# cp = 0 : minimum improvement in complexity parameter for splitting
fit = rpart(y ~ ., # formula
data=b.train, # dataframe used
control=rpart.control(xval=10, minsplitt=10, cp = 0.0))
nrow(fit$frame) # the biggest tree have 193 nodes
```

```
library(ROCR)
# start by predicting prob instead of class
y.pred = as.data.frame(predict(bal.fit.post, b.bal, type="prob"))
```

```
# first step is to compute the score object
y.pred.score = rediction(y.pred[,2], b.bal$y)
```

```
# now we will determine the optimal cutoff
# rather than use ROCR curve directly, the cost built into ROC is used
y.cost = performance(y.pred.score, measure="cost", cost.fn=5200, cost.fp=8700)
cutoff.best = y.cost@x.values[[1]][which.min(y.cost@y.values[[1]])] # 0.7142857
```

```
# make predictions using this cutoff rate for the test set
y.pred.test = predict(bal.fit.post, b.test, type="prob")
y.pred.test.cutoff = ifelse(y.pred.test[,2] > cutoff.best, 'yes', 'no')
```

```
# now find the profit for the test data set using the optimal cutoff
cm = confusionMatrix(table(pred=y.pred.test.cutoff,actual = b.test$y), positive='yes')
profit = 5200 - 5200 * (1-cm$byClass["Sensitivity"])[[1]] - 8700 * (1-cm$byClass["Specificity"])[[1]]
# profit = $2261.017
```

*See complete code on <https://github.com/Shelly-Yang/Predictive-Modeling-for-Direct-Marketing>.



Contents

Executive Summary

Analytics Goal and Approach

Appendix

Dataset Description

Feature	Description
age	Age of customer in years (integer)
job	Job category, one of 12 possible values, (string)
marital	Marital status, one of 3 possible values, (string)
education	Level of education, one of 4 possible values, (string)
default	Whether the customer has previously defaulted on a loan (yes or no)
balance	Average daily balance (numeric)
housing	Whether the customer has a mortgage with the bank (yes or no)
loan	Whether the customer has a revolving loan credit with the bank (yes or no)
day	Date of last contact
month	Month of last contact
contact	Method of last contact, one of three values (string)
duration	Duration of last contact, in minutes, (integer)
campaign	Campaign identifier
pdays	Days to prior contact
previous	Number of previous contacts
poutcome	Outcome of previous
y	Whether the customer purchased a CD or not (yes, or no), the class variable